

Random Forests over normalized data in CPU-GPU DBMSes

Zachary Huang, Pavan Kalyan Damalapati, Rathijit Sen, Eugene Wu
Data Science Institute, Columbia University, Microsoft Gray Systems Lab



COLUMBIA
UNIVERSITY



Microsoft

ML in databases is popular



Microsoft Azure

Using Azure Machine Learning with SQL Data Warehouse

Posted on November 24, 2015



[Sahaj Saini](#)

Program Manager, SQL Engineering

Azure Machine Learning allows you to build predictive models using data from your Azure SQL Data Warehouse database and other sources.

Azure Machine Learning is a powerful cloud-based predictive analytics service that makes it possible to



Google Cloud

What is BigQuery ML?

BigQuery ML lets you create and execute machine learning models using GoogleSQL queries. BigQuery ML democratizes machine learning by letting SQL practitioners build models using existing SQL tools and skills. BigQuery ML increases development speed by eliminating the need to move data.

BigQuery ML functionality is available



Amazon Redshift ML

Create, train, and deploy machine learning (ML) models using familiar SQL commands


[Read the overview guide](#)

Amazon Redshift ML makes it easy for data analysts and database developers to create, train, and apply machine learning models using familiar SQL commands in Amazon Redshift data warehouses. With Redshift ML, you can take advantage of [Amazon SageMaker](#) a

Random forests

Kaggle Survey 2022: All Results

Most popular machine learning algorithms in 2022



GPU for Random Forests

 **NVIDIA**
DEVELOPER

Join 

Data Science

Accelerating Random Forests Up to 45x Using cuML

Feb 25, 2021  0 Like  Discuss (1)

By [Vishal Mehta](#)



LightGBM GPU Tutorial

The purpose of this document is to give you a quick step-by-step tutorial on GPU training.

For Windows, please see [GPU Windows Tutorial](#).

We will use the GPU instance on [Microsoft Azure cloud computing platform](#) for demonstration, but you can use any machine with modern AMD or NVIDIA GPUs.



XGBoost GPU Support

This page contains information about GPU algorithms supported in XGBoost.


Note

CUDA 11.0, Compute Capability 5.0 required (See [this list](#) to look up compute capability of your GPU card.)

CUDA Accelerated Tree Construction Algorithms

Most of the algorithms in XGBoost including training, prediction and evaluation can be accelerated with CUDA-capable GPUs.

Challenge: DB-ML mismatch



DB: stores **many** tables

City	Type	Season	Name	Revenue
...

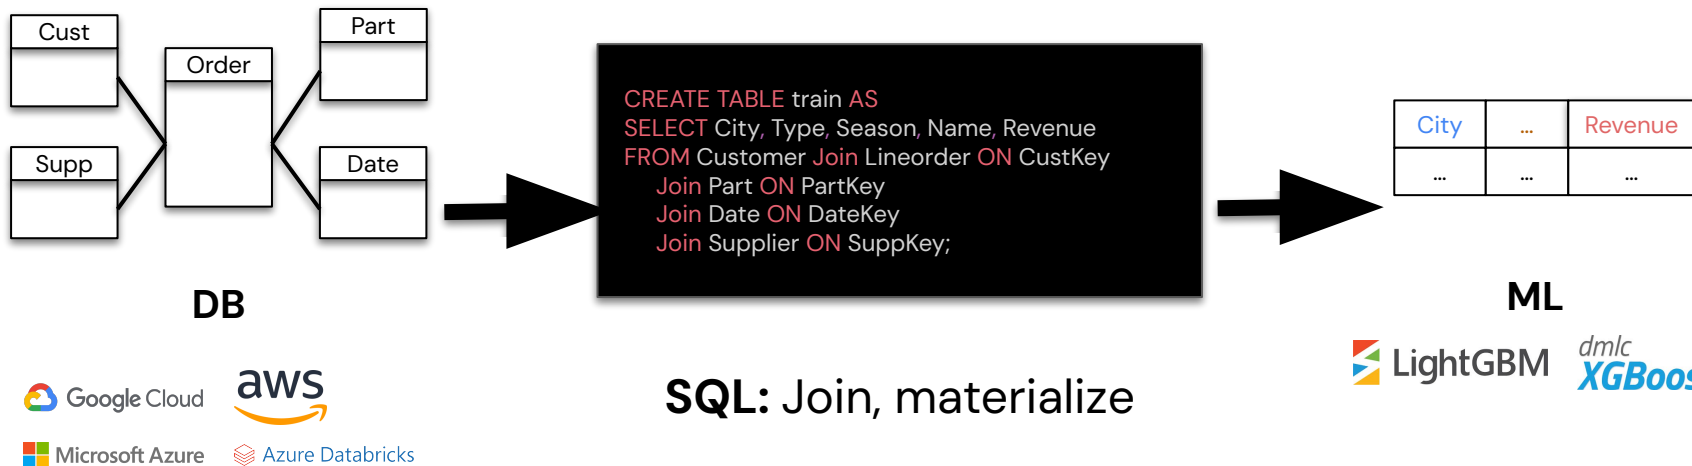
feature X y

ML: needs **single** table




Challenge: DB-ML mismatch

Current solution



Challenge: Join Results Exceed GPU Memory

- Join significantly increases size
- Common GPU memory is limited



↓ **Join size >15 GB**

City	...	Revenue
...

GPU	memory
V100	16/32 GB
P100	16 GB
K80	12 GB
P4	8 GB

In contrast, CPUs can easily support 100GB to several TBs memory

Our Approach

Algorithm:

- Translate Random Forests Training as Join Aggregations
- Push down Aggregation before Join

System:

- Explore design choices on CPU-GPU DBMSes

We will start with the background of Random Forests Training

Background: Random Forests Training

Feature Target

 R

F	Y
1	1
1	1
2	1
2	2
3	3
3	4

Random Forests are ensemble of trees. For each tree:

Tree-split Problem: Find the best split point F to minimize the sum of variance of Y.

$$\text{var}([1,1,1,2,3,4]) \\ = 1.6$$

Background: Random Forests Training

R

F	Y
1	1
1	1
2	1
2	2
3	3
3	4

Tree-split Problem: Find the best split point F to minimize the sum of variance of Y.

$$F > 1$$

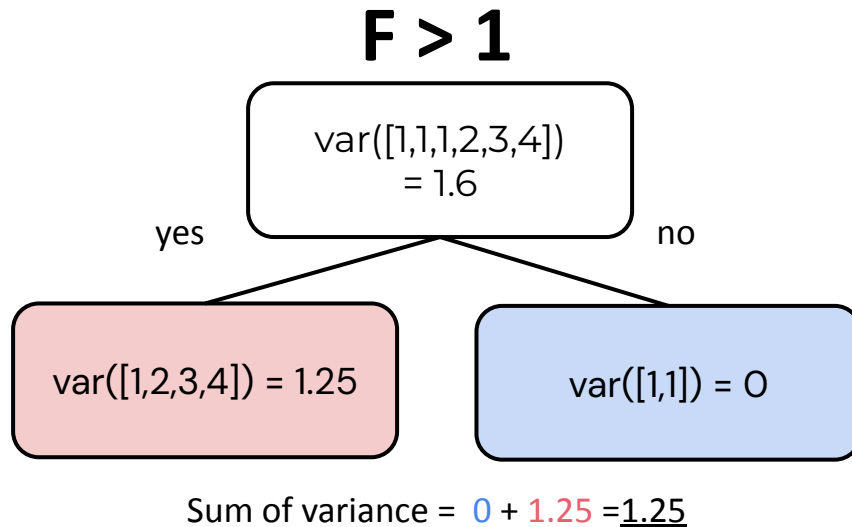
$$\text{var}([1,1,1,2,3,4]) \\ = 1.6$$

Background: Random Forests Training

R

F	Y
1	1
1	1
2	1
2	2
3	3
3	4

Tree-split Problem: Find the best split point F to minimize the sum of variance of Y.

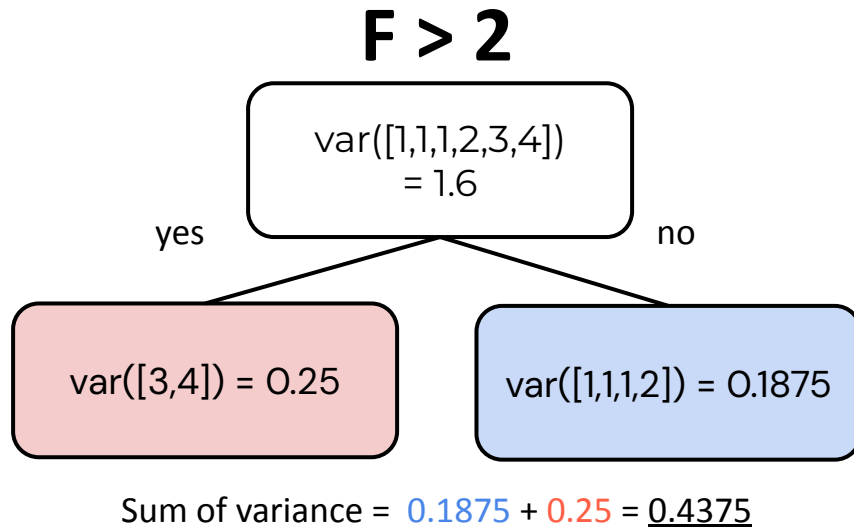


Background: Random Forests Training

R

F	Y
1	1
1	1
2	1
2	2
3	3
3	4

Tree-split Problem: Find the best split point F to minimize the sum of variance of Y.

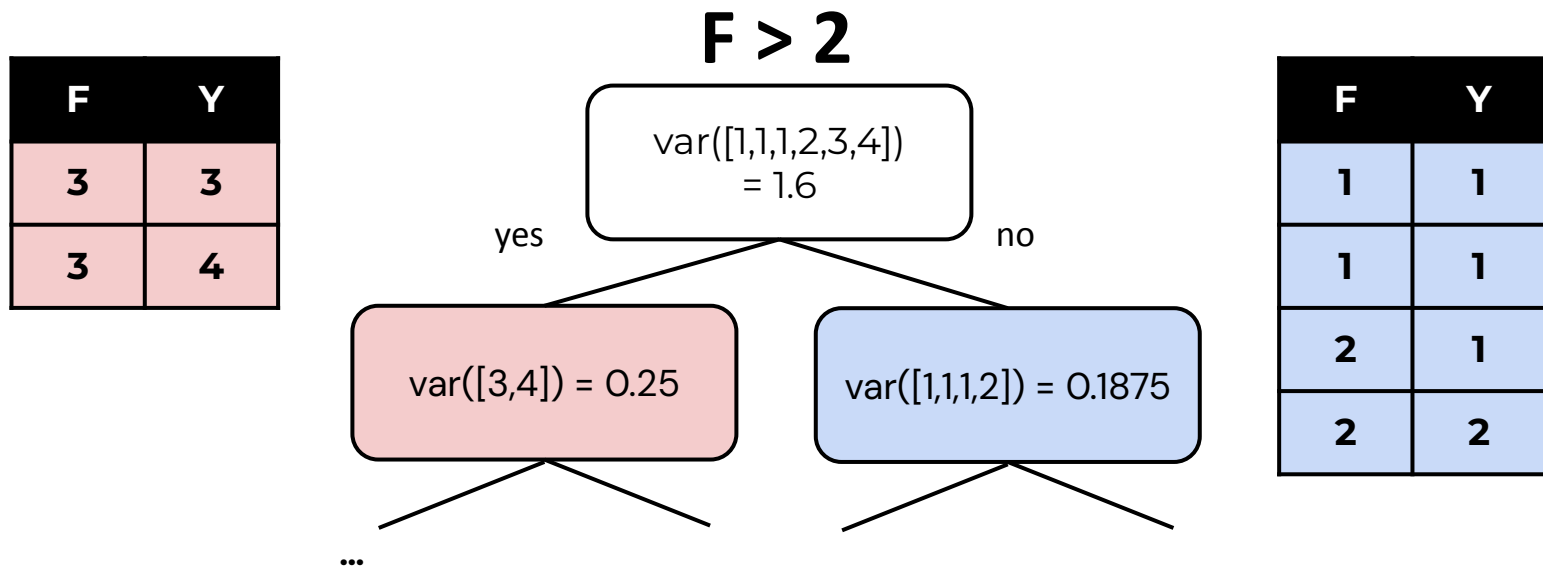


$F > 2$ is the best split point!

Background: Random Forests Training

We have splitted the base relation.

Recursively split the splitted relation and find the best splits.



Core Computation: Variance

$$\text{Variance}(Y) = \Sigma Y^2 - (\Sigma Y)^2 / \Sigma 1$$

Based on sum of squares, sum, count

Lightweight Postprocessing
To find the feature with the best split

```
SELECT F, - S * S / C AS variance  
FROM (
```

```
SELECT F, SUM(1) AS C, SUM(Y) AS S, SUM(Y * Y) AS Q  
FROM R  
GROUP BY F
```

Core Aggregation

```
)  
GROUP BY F  
ORDER BY variance DESC LIMIT 1;
```

*The query for tree is more complex by computing the **sum** of variances. Please refer to the paper for the detailed queries.

Core Computation: Variance

$$\text{Variance}(Y) = \Sigma Y^2 - (\Sigma Y)^2 / \Sigma 1$$

Based on sum of squares, sum, count

```
SELECT F, SUM(1) AS C, SUM(Y) AS S, SUM(Y * Y) AS Q  
FROM R  
GROUP BY F
```

Core Aggregation



*The query for tree is more complex by computing the **sum** of variances. Please refer to the paper for the detailed queries.

Core Computation: Variance

$$\text{Variance}(Y) = \Sigma Y^2 - (\Sigma Y)^2 / \Sigma 1$$

Based on sum of squares, sum, count

How to speed up Core Aggregation computation over join?

```
SELECT F, SUM(1) AS C, SUM(Y) AS S, SUM(Y * Y) AS Q  
FROM R JOIN F ON J  
GROUP BY F
```



*The query for tree is more complex by computing the **sum** of variances. Please refer to the paper for the detailed queries.

Our Solution: Aggregation Pushdown

```
SELECT F, SUM(1) AS C, SUM(Y) AS S, SUM(Y * Y) AS Q  
FROM R JOIN F ON J  
GROUP BY F
```

F (fact)

Y	J
1	1
2	1
3	2
4	3


R (dim)

F	J
1	1
1	2
2	3

Our Solution: Aggregation Pushdown

```
SELECT F, SUM(1) AS C, SUM(Y) AS S, SUM(Y * Y) AS Q  
FROM R JOIN F ON J  
GROUP BY F
```

As large as the fact table



F (fact)

Y	J
1	1
2	1
3	2
4	3

R (dim)

F	J
1	1
1	2
2	3

$R \bowtie F$

Y	J	F
1	1	1
2	1	1
3	2	1
4	3	2

Our Solution: Aggregation Pushdown

```
SELECT F, SUM(1) AS C, SUM(Y) AS S, SUM(Y * Y) AS Q  
FROM R JOIN F ON J  
GROUP BY F
```

F (fact)

Y	J
1	1
2	1
3	2
4	3

R (dim)

F	J
1	1
1	2
2	3

Our Solution: Aggregation Pushdown

```
CREATE TABLE msg AS  
SELECT J, SUM(1) AS C, SUM(Y) AS S, SUM(Y * Y) AS Q  
FROM F GROUP BY J;
```

F (fact)

Y	J
1	1
2	1
3	2
4	3

Msg (partial agg)

J	(C,S,Q)
1	(2,3,5)
2	(1,3,9)
3	(1,4,16)



R (dim)

F	J
1	1
1	2
2	3

Our Solution: Aggregation Pushdown

```
SELECT F, SUM(C) AS C, SUM(S) AS S, SUM(Q) AS Q  
FROM R JOIN msg ON J  
GROUP BY F
```

Bounded by
dimension table

F (fact)

Y	J
1	1
2	1
3	2
4	3

Msg (partial agg)

J	(C,S,Q)
1	(2,3,5)
2	(1,3,9)
3	(1,4,16)

R (dim)

F	J
1	1
1	2
2	3

R ⋈ Msg

F	(C,S,Q)
1	(3,6,14)
2	(1,4,16)

Our Solution: Aggregation Pushdown

```
SELECT F, SUM(1) AS C, SUM(Y) AS S, SUM(Y * Y) AS Q  
FROM R JOIN F ON J  
GROUP BY F
```

Aggregation over Join



Query rewrite

```
CREATE TABLE msg AS  
SELECT SUM(1) AS C, SUM(Y) AS S, SUM(Y * Y) AS Q  
FROM F GROUP BY J;
```

Partial Aggregate
from Fact Table

```
SELECT F, SUM(C) AS C, SUM(S) AS S, SUM(Q) AS Q  
FROM R JOIN msg ON J  
GROUP BY F
```

Sum Partial Aggregate
by Dim Table

System Design

Problem: Train random forests directly in CPU-GPU DBMSes.
We can use CPU memory when data doesn't fit in GPU.

Where to place data and execute query for CPU & GPU?

Aggregation for random forests

Msg from Fact table

Simple query over big data
Better on GPU

```
CREATE TABLE msg AS  
SELECT SUM(1) AS C, SUM(Y) AS S, SUM(Y * Y) AS Q  
FROM F GROUP BY J;
```

Lightweight Postprocessing by Dim table

Complex query over small data
Better on CPU
Inter-query Parallelism

```
SELECT F, Q - S * S / C AS variance  
FROM (  
  
    SELECT F, SUM(C) AS C, SUM(S) AS S, SUM(Q) AS Q  
    FROM R JOIN msg ON J  
    GROUP BY F  
  
)  
GROUP BY F  
GROUP BY variance DESC LIMIT 1;
```


Aggregation for random forests



Msg from Fact table

Simple query over big data
Better on GPU

```
CREATE TABLE msg AS  
SELECT SUM(1) AS C, SUM(Y) AS S, SUM(Y * Y) AS Q  
FROM F GROUP BY J;
```

What if F doesn't fit in GPU memory?

Store only a horizontal partition of F fit in GPU.




Aggregation for random forests

Msg from Fact table

Simple query over big data
Better on GPU

```
CREATE TABLE msg AS  
SELECT SUM(1) AS C, SUM(Y) AS S, SUM(Y * Y) AS Q  
FROM F GROUP BY J;
```


How to execute query over CPU part of F?




Aggregation for random forests

How to execute query over CPU part of F?

1. Directly use the CPU to execute query



2. Transfer through PCIe from CPU to GPU, then execute query using GPU



Implementation & Experiment



We use off-the-shelf dataframes to integrate with the Python ecosystem.
Aggregation pushdown achieved through query rewriting.

	Data	Query Execution
CPU	Pandas DF	DuckDB
GPU	CuDF	CuDF

*Pandas DF is not efficient (single thread), so use DuckDB.

* cuDF is not efficient as it materializes intermediates, and our results can be improved using backends like CUDA based on Crystal

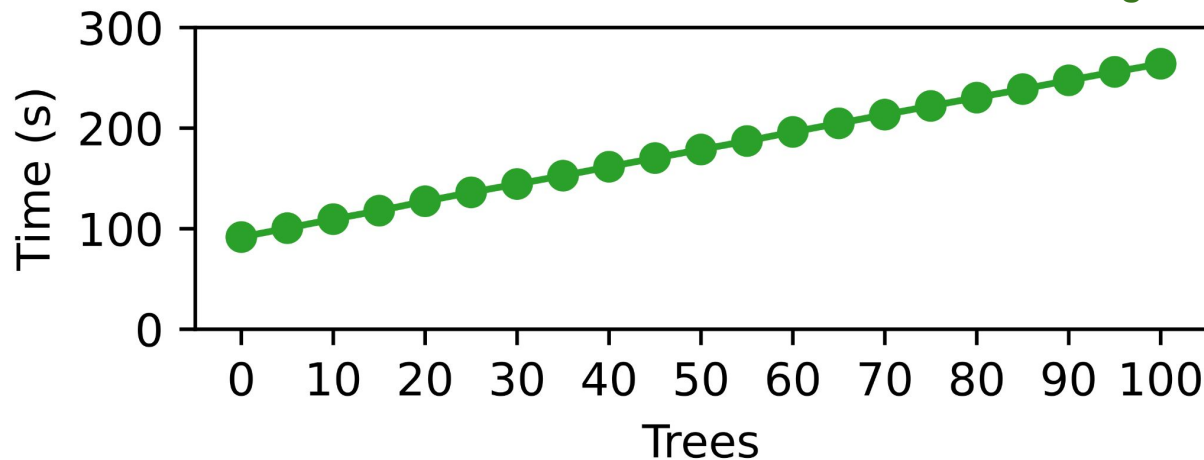
Hardware: V100 GPU (16 GB) connected via PCIe3 to CPU (112 GB)

Task: Random forests (100 trees, 8 max leaves, 0.5 sample) for SSB dataset

When the database fit in GPU


SSB SF=10 (6.5GB after encoding)

LightGBM requires Join



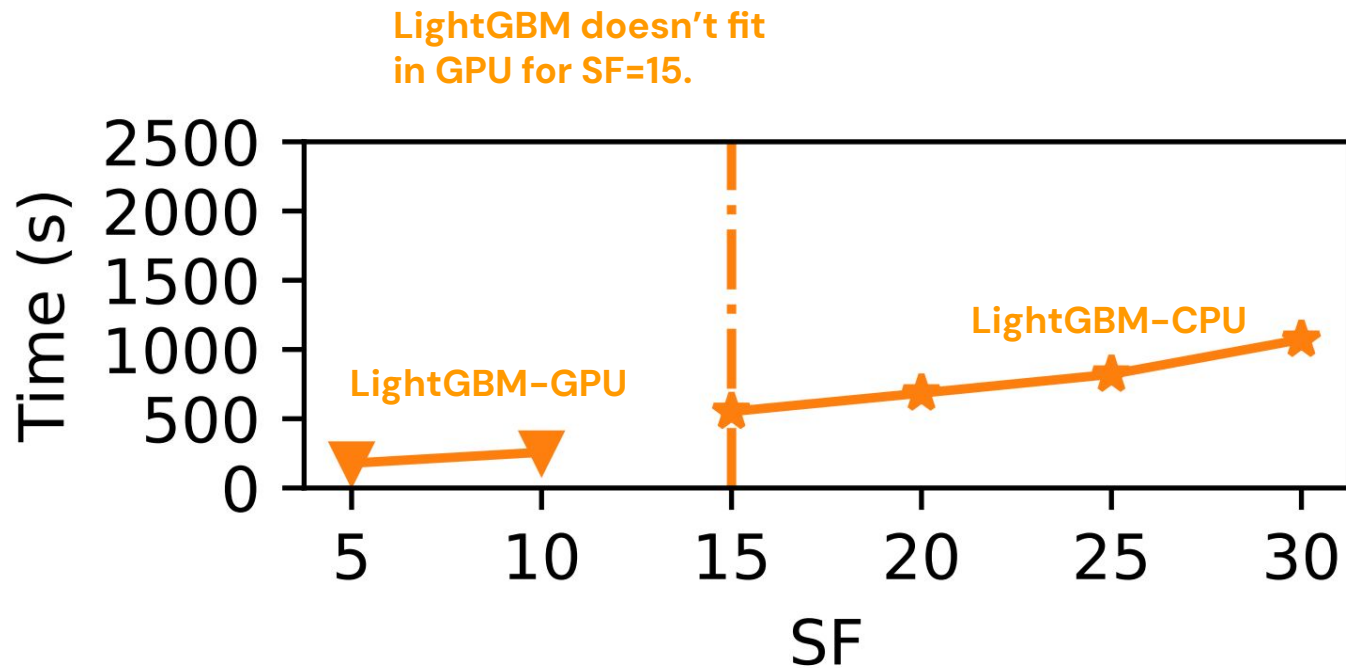
When the database fit in GPU

SSB SF=10 (6.5GB after encoding)



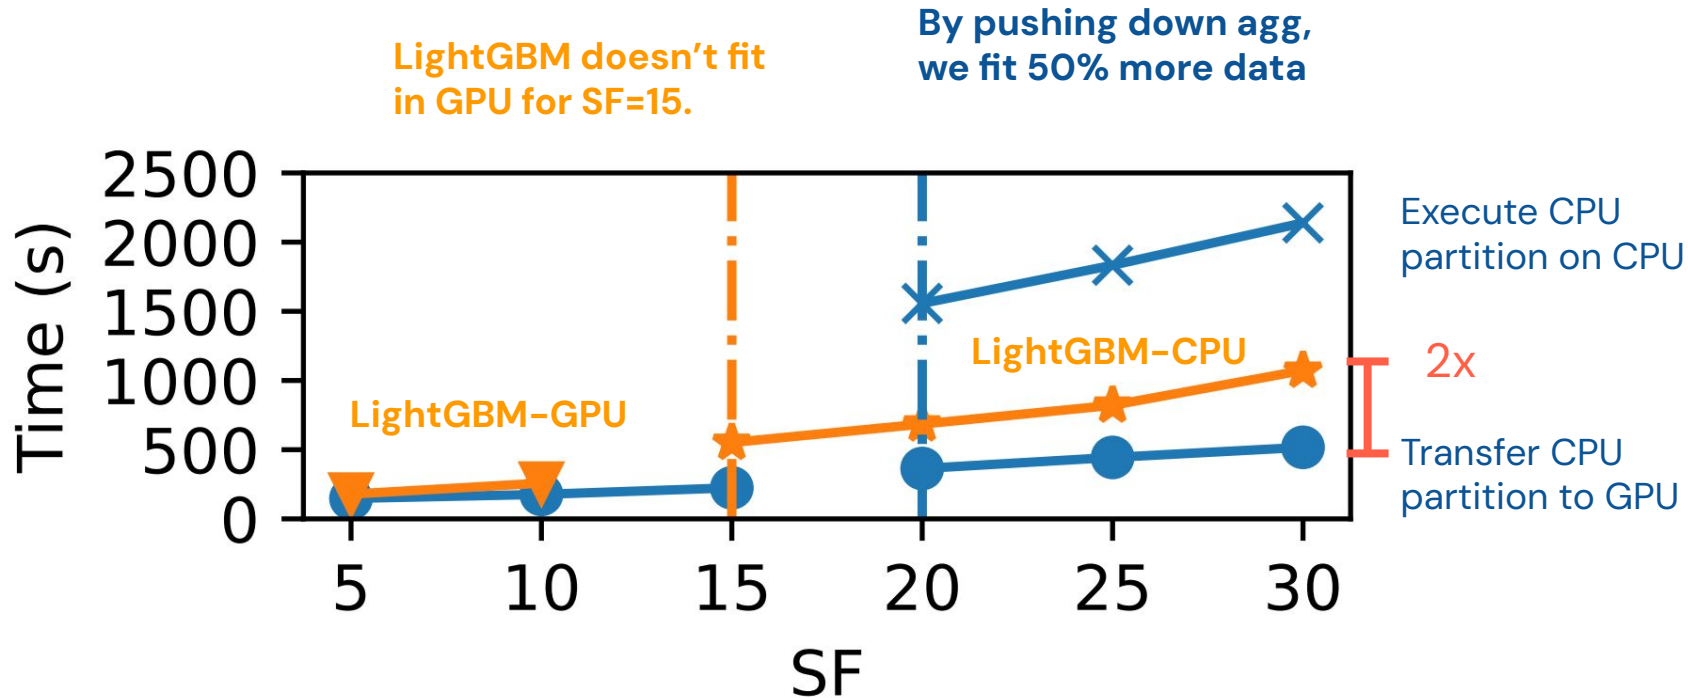
When the database doesn't fit in GPU

Vary SSB SF



When the database doesn't fit in GPU

Vary SSB SF



Thank you!

Repo



Demo

